

# Markdown\_fa.R

yaxin

2021-04-12

```
# Uncomment if packages not installed
## install.packages("psych")
## install.packages("caret")
## install.packages("randomForest")
## install.packages("MLmetrics")
## install.packages("doParallel")
## install.packages("kernlab")
## install.packages("glmnet")

# Load data
setwd('D:\\Yaxin\\HKBU BM\\Courses\\Sem 2\\ECON7860 Big Data Analytics for Business (S11)\\Group Project')
rawData <- read.csv2("HR_comma_sep.csv", sep = ',')
colnames(rawData)

## [1] "satisfaction_level"      "last_evaluation"        "number_project"
## [4] "average_montly_hours"    "time_spend_company"     "Work_accident"
## [7] "left"                    "promotion_last_5years"  "sales"
## [10] "salary"

# Move the target variable "left" after "time_spend_company"
library(dplyr)

## Warning: package 'dplyr' was built under R version 4.0.4

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```

rawData <- rawData %>% relocate(left, .after = time_spend_company)

# Transform feature types
transform_feature <- function(X) {
  X$satisfaction_level <- as.numeric(X$satisfaction_level)
  X$last_evaluation <- as.numeric(X$last_evaluation)
  X$Work_accident <- as.factor(X$Work_accident)
  X$promotion_last_5years <- as.factor(X$promotion_last_5years)
  X$sales <- as.factor(X$sales)
  X$salary <- as.factor(X$salary)
  X$left <- factor(ifelse(X$left == 0, 'no', 'yes'), levels = c('yes', 'no'))
  return(X)
}

rawData <- transform_feature(rawData)
summary(rawData)

```

```

##  satisfaction_level last_evaluation  number_project  average_monthly_hours
##  Min.      :0.0900    Min.      :0.3600    Min.      :2.000    Min.      : 96.0
##  1st Qu.:0.4400    1st Qu.:0.5600    1st Qu.:3.000    1st Qu.:156.0
##  Median :0.6400    Median :0.7200    Median :4.000    Median :200.0
##  Mean   :0.6128    Mean   :0.7161    Mean   :3.803    Mean   :201.1
##  3rd Qu.:0.8200    3rd Qu.:0.8700    3rd Qu.:5.000    3rd Qu.:245.0
##  Max.   :1.0000    Max.   :1.0000    Max.   :7.000    Max.   :310.0
##
##  time_spend_company  left      Work_accident  promotion_last_5years
##  Min.      : 2.000    yes: 3571    0:12830    0:14680
##  1st Qu.: 3.000    no :11428    1: 2169    1: 319
##  Median : 3.000
##  Mean   : 3.498
##  3rd Qu.: 4.000
##  Max.   :10.000
##
##           sales      salary
##  sales      :4140    high :1237
##  technical   :2720    low  :7316
##  support     :2229    medium:6446
##  IT          :1227
##  product_mng: 902
##  marketing   : 858
##  (Other)     :2923

```

```

# Separate features and target variable
X <- rawData
y <- X$left
tag <- colnames(X)
tag

```

```

##  [1] "satisfaction_level"  "last_evaluation"    "number_project"
##  [4] "average_monthly_hours" "time_spend_company" "left"

```

```
## [7] "Work_accident"      "promotion_last_5years" "sales"
## [10] "salary"
```

```
# Feature engineering
## Define the function for factor extraction
require(psych)
```

```
## Loading required package: psych
```

```
## Warning: package 'psych' was built under R version 4.0.4
```

```
fe <- function(M, n) {
  # The numeric feature matrix M needs to be normalized beforehand
  fa1 <- fa(M, n)
  fa.diagram(fa1)
  return(list('scores' = fa1$scores, 'weights' = fa1$weights))
}
```

```
## Create dummy variables for "sales" and "salary"
dummySales <- dummy.code(X$sales)
dummySalary <- dummy.code(X$salary)
colnames(dummySales)
```

```
## [1] "sales"      "technical"  "support"    "IT"         "product_mng"
## [6] "marketing"  "RandD"      "accounting" "hr"         "management"
```

```
colnames(dummySalary)
```

```
## [1] "low"      "medium" "high"
```

```
### Set "sales" and "low" as the default values respectively
dummySales <- dummySales[, -c(1)]
dummySalary <- dummySalary[, -c(1)]
```

```
X_dummy <- cbind(X[, -c(9, 10)], dummySales, dummySalary)
tag_dummy <- colnames(X_dummy)
tag_dummy
```

```
## [1] "satisfaction_level"  "last_evaluation"    "number_project"
## [4] "average_monthly_hours" "time_spend_company" "left"
## [7] "Work_accident"      "promotion_last_5years" "technical"
## [10] "support"            "IT"                 "product_mng"
## [13] "marketing"          "RandD"              "accounting"
## [16] "hr"                 "management"         "medium"
## [19] "high"
```

```
# Train(80%)-test(20%)-split (stratified as "left" is unbalanced)
library(caret)
```

```

## Warning: package 'caret' was built under R version 4.0.4

## Loading required package: lattice

## Loading required package: ggplot2

##
## Attaching package: 'ggplot2'

## The following objects are masked from 'package:psych':
##
##      %+%, alpha

## Set seed for replication purpose
set.seed(7860)
index <- createDataPartition(y, p = 0.8, list = FALSE)
X_train <- X[index, ]
X_test <- cbind(X[-index, 1 : 5], X[-index, 7 : length(X)])
X_dummy_train <- X_dummy[index, ]
X_dummy_test <- cbind(X_dummy[-index, 1 : 5], X_dummy[-index, 7 : length(X_dummy)])
y_test <- y[-index]

# Modeling with extracted factors, 5-fold nested CV with random search
n_fold <- 5
cv <- createFolds(X_train$left, n_fold)
#models <- c('glmnet', 'sumLinear', 'rf', 'knn')
models <- c('glmnet', 'rf', 'knn')
tune <- 15
n_cluster <- 5 ## Please set the number of multiprocessing slaves accordingly

for (m in models) {
  assign(paste0(m, '_cv'),
    list('1' = list('model' = c(), 'f1_val' = c(), 'confm' = c()),
        '2' = list('model' = c(), 'f1_val' = c(), 'confm' = c()),
        '3' = list('model' = c(), 'f1_val' = c(), 'confm' = c()),
        '4' = list('model' = c(), 'f1_val' = c(), 'confm' = c()),
        '5' = list('model' = c(), 'f1_val' = c(), 'confm' = c()))

  for (i in 1 : n_fold) {
    ## Extract factors and weights from EFA
    fa_result <- fe(scale(X_train[-cv[[i]], 1 : 5]), 2)
    performance <- fa_result$scores[ , 1]
    satisfaction <- fa_result$scores[ , 2]
    w <- fa_result$weights
    colnames(w) <- c('performance', 'satisfaction')
    y_val <- X_train[cv[[i]], 'left']

    control <- trainControl(method = 'repeatedcv', number = 5, repeats = 2,
                           summaryFunction = prSummary, classProbs = TRUE,
                           search="random", verboseIter = TRUE)
  }
}

```

```

set.seed(7860)

require(doParallel)
cl <- makePSOCKcluster(n_cluster, outfile = '')
registerDoParallel(cl)

if (m == 'rf') {
  X_dev <- cbind(performance, satisfaction,
                 X_train[-cv[[i]], 6 : length(X_train)])
  X_val <- cbind(scale(as.matrix(X_train[cv[[i]], 1 : 5]) %*% w),
                 X_train[cv[[i]], 7 : length(X_train)])

  m1 <- train(left ~ ., data = X_dev, method = m, metric = 'F',
              tuneLength = tune, trControl = control)

  rf_cv[[i]][['model']] <- m1
  rf_cv[[i]][['f1_val']] <- F_meas(predict(m1, X_val), y_val)
  rf_cv[[i]][['confm']] <- confusionMatrix(predict(m1, X_val), y_val)
} else {
  X_dev <- cbind(performance, satisfaction,
                 X_dummy_train[-cv[[i]], 6 : length(X_dummy_train)])
  X_val <- cbind(scale(as.matrix(X_dummy_train[cv[[i]], 1 : 5]) %*% w),
                 X_dummy_train[cv[[i]], 7 : length(X_dummy_train)])

  if (m == 'glmnet') {
    m1 <- train(left ~ ., data = X_dev, method = m, family = 'binomial',
                metric = 'F', tuneLength = tune, trControl = control)

    glmnet_cv[[i]][['model']] <- m1
    glmnet_cv[[i]][['f1_val']] <- F_meas(predict(m1, X_val), y_val)
    glmnet_cv[[i]][['confm']] <- confusionMatrix(predict(m1, X_val), y_val)
  } else if (m == 'knn') {
    m1 <- train(left ~ ., data = X_dev, method = m, metric = 'F',
                tuneLength = tune, trControl = control,
                tuneGrid = expand.grid(k = c(2, 3, 4, 5, 10)))

    knn_cv[[i]][['model']] <- m1
    knn_cv[[i]][['f1_val']] <- F_meas(predict(m1, X_val), y_val)
    knn_cv[[i]][['confm']] <- confusionMatrix(predict(m1, X_val), y_val)
  } else if (m == 'svmLinear') {
    m1 <- train(left ~ ., data = X_dev, method = m, metric = 'F',
                tuneLength = tune, trControl = control)

    svmLinear_cv[[i]][['model']] <- m1
    svmLinear_cv[[i]][['f1_val']] <- F_meas(predict(m1, X_val), y_val)
    svmLinear_cv[[i]][['confm']] <- confusionMatrix(predict(m1, X_val), y_val)
  }
}

stopImplicitCluster()

```

```

    stopCluster(cl)
  }
}

## Loading required namespace: GPArotation

## Loading required package: doParallel

## Warning: package 'doParallel' was built under R version 4.0.4

## Loading required package: foreach

## Warning: package 'foreach' was built under R version 4.0.4

## Loading required package: iterators

## Warning: package 'iterators' was built under R version 4.0.4

## Loading required package: parallel

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results

## Selecting tuning parameters
## Fitting alpha = 0.153, lambda = 0.00124 on full training set

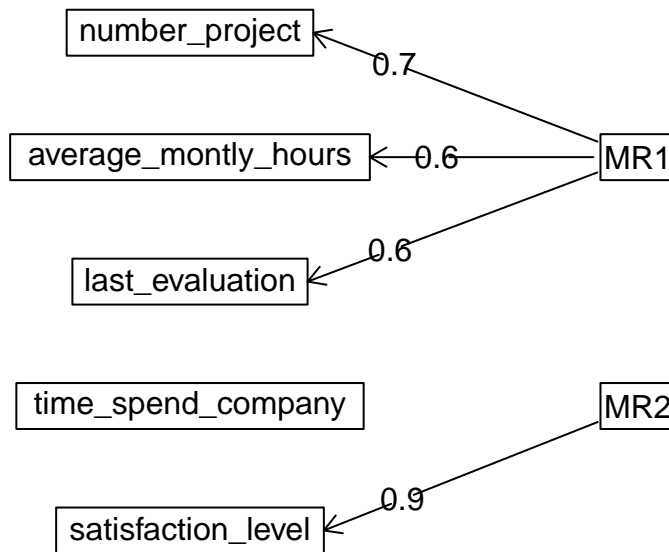
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results

```

## Factor Analysis



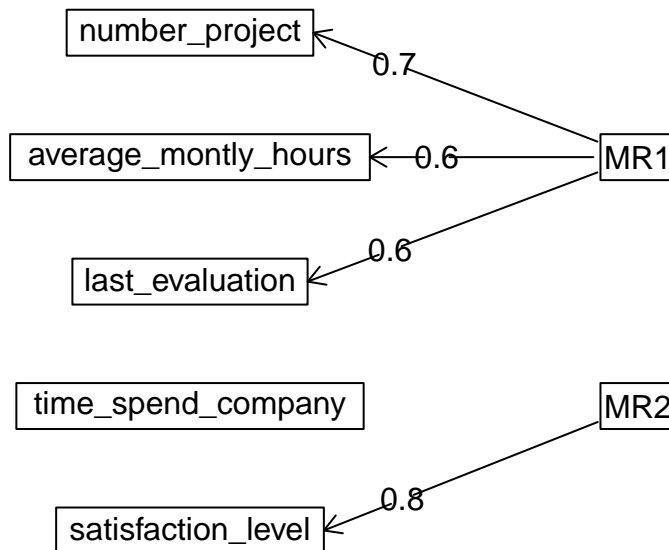
```
## Selecting tuning parameters
## Fitting alpha = 0.153, lambda = 0.00124 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results
```

## Factor Analysis



```
## Selecting tuning parameters
## Fitting alpha = 0.153, lambda = 0.00124 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results

## Selecting tuning parameters
## Fitting alpha = 0.153, lambda = 0.00124 on full training set


## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results
```



## Factor Analysis



```
## Selecting tuning parameters
## Fitting alpha = 0.572, lambda = 0.042 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.


## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 10 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 10 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

## Factor Analysis



```
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 10 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.


## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 8 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results
```

## Factor Analysis

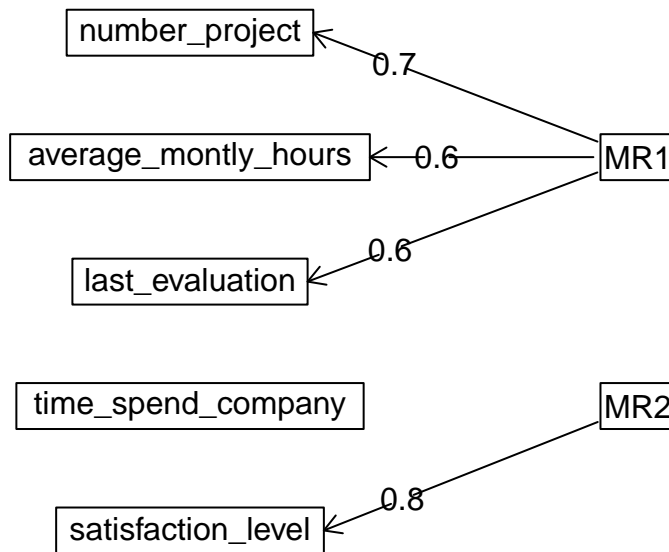


```
## Selecting tuning parameters  
## Fitting mtry = 12 on full training set
```

```
## Aggregating results  
## Selecting tuning parameters  
## Fitting k = 3 on full training set
```

```
## Aggregating results  
## Selecting tuning parameters  
## Fitting k = 3 on full training set
```


## Factor Analysis



```
## Aggregating results  
## Selecting tuning parameters  
## Fitting k = 3 on full training set
```

```
## Aggregating results  
## Selecting tuning parameters  
## Fitting k = 3 on full training set
```

## Factor Analysis



```
## Aggregating results
## Selecting tuning parameters
## Fitting k = 5 on full training set
```

```
# Print CV results
for (m in models) {
  cat(rep('\n', 3))
  print(get(paste0(m, '_cv')))
  cat(rep('\n', 3))
}
```

```
##
##
##
## $'1'
## $'1'$model
## glmnet
##
## 9601 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7681, 7681, 7681, 7681, 7681, ...
```

```

## Resampling results across tuning parameters:
##
##   alpha      lambda      AUC      Precision  Recall      F
##   0.03431740 0.089848573 0.5772080 0.7876377 0.1854710 0.2998691
##   0.08249613 0.360573416 0.5833627      NaN 0.0000000      NaN
##   0.11032045 5.835644047 0.0000000      NaN 0.0000000      NaN
##   0.14298028 5.900723701 0.0000000      NaN 0.0000000      NaN
##   0.15333117 0.001237096 0.5727423 0.5898577 0.2865355 0.3853157
##   0.24414545 0.041843375 0.5832865 0.6731357 0.2303149 0.3429387
##   0.24717296 0.005363840 0.5746734 0.5736787 0.2567858 0.3545295
##   0.29186267 0.038180467 0.5838735 0.6618070 0.2318462 0.3432183
##   0.37615343 0.002088687 0.5737165 0.5849002 0.2786604 0.3771301
##   0.57222747 0.041985838 0.5869645 0.7086693 0.2318457 0.3491847
##   0.73301844 2.449870885 0.0000000      NaN 0.0000000      NaN
##   0.77100396 0.191331587 0.5736713      NaN 0.0000000      NaN
##   0.78001638 0.014048926 0.5836553 0.5934982 0.2392822 0.3409883
##   0.96830053 0.005986254 0.5790393 0.5678478 0.2440952 0.3413153
##   0.98979898 1.565802338 0.0000000      NaN 0.0000000      NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1533312 and lambda
## = 0.001237096.
##
## $'1'$f1_val
## [1] 0.4175084
##
## $'1'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes 186 134
##           no  385 1694
##
##           Accuracy : 0.7837
##           95% CI : (0.7666, 0.8)
##           No Information Rate : 0.762
##           P-Value [Acc > NIR] : 0.006362
##
##           Kappa : 0.2974
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.32574
##           Specificity : 0.92670
##           Pos Pred Value : 0.58125
##           Neg Pred Value : 0.81481
##           Prevalence : 0.23802
##           Detection Rate : 0.07753
##           Detection Prevalence : 0.13339
##           Balanced Accuracy : 0.62622
##
##           'Positive' Class : yes
##

```

```

##
##
## $'2'
## $'2'$model
## glmnet
##
## 9600 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7680, 7680, 7680, 7680, 7680, ...
## Resampling results across tuning parameters:
##
##   alpha      lambda      AUC      Precision  Recall      F
##   0.03431740 0.089848573 0.5807717 0.8006113 0.2107221 0.3332538
##   0.08249613 0.360573416 0.5792088      NaN 0.0000000      NaN
##   0.11032045 5.835644047 0.0000000      NaN 0.0000000      NaN
##   0.14298028 5.900723701 0.0000000      NaN 0.0000000      NaN
##   0.15333117 0.001237096 0.5761376 0.5941547 0.2820569 0.3821086
##   0.24414545 0.041843375 0.5841563 0.6925909 0.2413567 0.3576818
##   0.24717296 0.005363840 0.5780074 0.5851413 0.2621444 0.3617453
##   0.29186267 0.038180467 0.5844313 0.6814771 0.2415755 0.3564603
##   0.37615343 0.002088687 0.5770560 0.5890117 0.2746171 0.3741422
##   0.57222747 0.041985838 0.5843376 0.7184135 0.2413567 0.3610837
##   0.73301844 2.449870885 0.0000000      NaN 0.0000000      NaN
##   0.77100396 0.191331587 0.5660670      NaN 0.0000000      NaN
##   0.78001638 0.014048926 0.5844165 0.6008996 0.2461707 0.3490103
##   0.96830053 0.005986254 0.5812564 0.5764578 0.2503282 0.3488117
##   0.98979898 1.565802338 0.0000000      NaN 0.0000000      NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1533312 and lambda
## = 0.001237096.
##
## $'2'$f1_val
## [1] 0.372549
##
## $'2'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction yes  no
##      yes  171  175
##      no   401 1653
##
##           Accuracy : 0.76
##           95% CI : (0.7424, 0.777)
##      No Information Rate : 0.7617
##      P-Value [Acc > NIR] : 0.5869
##
##           Kappa : 0.2351
##

```

```

## McNemar's Test P-Value : <2e-16
##
##      Sensitivity : 0.29895
##      Specificity : 0.90427
##      Pos Pred Value : 0.49422
##      Neg Pred Value : 0.80477
##      Prevalence : 0.23833
##      Detection Rate : 0.07125
##      Detection Prevalence : 0.14417
##      Balanced Accuracy : 0.60161
##
##      'Positive' Class : yes
##
##
##
## '$3'
## '$3'$model
## glmnet
##
## 9600 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
##      alpha      lambda      AUC      Precision  Recall      F
## 0.03431740 0.089848573 0.5839719 0.8064638 0.2174085 0.3417013
## 0.08249613 0.360573416 0.5843289      NaN 0.0000000      NaN
## 0.11032045 5.835644047 0.0000000      NaN 0.0000000      NaN
## 0.14298028 5.900723701 0.0000000      NaN 0.0000000      NaN
## 0.15333117 0.001237096 0.5796851 0.5938845 0.2882631 0.3872137
## 0.24414545 0.041843375 0.5866031 0.6882912 0.2482461 0.3639817
## 0.24717296 0.005363840 0.5812199 0.5852592 0.2679273 0.3668093
## 0.29186267 0.038180467 0.5868749 0.6769617 0.2497778 0.3639797
## 0.37615343 0.002088687 0.5803097 0.5906427 0.2830143 0.3818531
## 0.57222747 0.041985838 0.5878423 0.7165020 0.2471539 0.3665754
## 0.73301844 2.449870885 0.0000000      NaN 0.0000000      NaN
## 0.77100396 0.191331587 0.5695331      NaN 0.0000000      NaN
## 0.78001638 0.014048926 0.5867639 0.6015270 0.2543697 0.3568058
## 0.96830053 0.005986254 0.5844679 0.5821560 0.2594006 0.3581493
## 0.98979898 1.565802338 0.0000000      NaN 0.0000000      NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1533312 and lambda
## = 0.001237096.
##
## '$3'$f1_val
## [1] 0.3917749
##
## '$3'$confm
## Confusion Matrix and Statistics

```



```

##
##           Reference
## Prediction  yes   no
##           yes  181 172
##           no   390 1657
##
##           Accuracy : 0.7658
##           95% CI : (0.7484, 0.7827)
##           No Information Rate : 0.7621
##           P-Value [Acc > NIR] : 0.3431
##
##           Kappa : 0.2566
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.31699
##           Specificity : 0.90596
##           Pos Pred Value : 0.51275
##           Neg Pred Value : 0.80948
##           Prevalence : 0.23792
##           Detection Rate : 0.07542
##           Detection Prevalence : 0.14708
##           Balanced Accuracy : 0.61147
##
##           'Positive' Class : yes
##
##
##
## '$'4'
## '$'4'$model
## glmnet
##
## 9600 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
##  alpha      lambda      AUC      Precision  Recall      F
##  0.03431740  0.089848573  0.5846583  0.7855777  0.2038523  0.3231638
##  0.08249613  0.360573416  0.5846747      NaN  0.0000000      NaN
##  0.11032045  5.835644047  0.0000000      NaN  0.0000000      NaN
##  0.14298028  5.900723701  0.0000000      NaN  0.0000000      NaN
##  0.15333117  0.001237096  0.5811456  0.5988584  0.2928717  0.3928077
##  0.24414545  0.041843375  0.5887475  0.6769320  0.2410275  0.3550853
##  0.24717296  0.005363840  0.5829132  0.5885299  0.2703396  0.3700639
##  0.29186267  0.038180467  0.5891996  0.6713757  0.2445276  0.3580867
##  0.37615343  0.002088687  0.5820901  0.5944076  0.2856535  0.3853012
##  0.57222747  0.041985838  0.5890402  0.7036780  0.2421225  0.3597007
##  0.73301844  2.449870885  0.0000000      NaN  0.0000000      NaN
##  0.77100396  0.191331587  0.5727108      NaN  0.0000000      NaN

```

```

## 0.78001638 0.014048926 0.5888988 0.6018243 0.2517458 0.3547284
## 0.96830053 0.005986254 0.5857254 0.5808015 0.2580901 0.3570937
## 0.98979898 1.565802338 0.0000000      NaN 0.0000000      NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1533312 and lambda
## = 0.001237096.
##
## $'4'$f1_val
## [1] 0.3471616
##
## $'4'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes  159  186
##           no   412 1643
##
##           Accuracy : 0.7508
##           95% CI : (0.733, 0.768)
##           No Information Rate : 0.7621
##           P-Value [Acc > NIR] : 0.9058
##
##           Kappa : 0.2046
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.27846
##           Specificity : 0.89831
##           Pos Pred Value : 0.46087
##           Neg Pred Value : 0.79951
##           Prevalence : 0.23792
##           Detection Rate : 0.06625
##           Detection Prevalence : 0.14375
##           Balanced Accuracy : 0.58838
##
##           'Positive' Class : yes
##
##
## $'5'
## $'5'$model
## glmnet
##
## 9599 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##

```

```

##      alpha      lambda      AUC      Precision      Recall      F
##  0.03431740  0.089848573  0.5870402  0.8210220  0.2284464  0.3570610
##  0.08249613  0.360573416  0.5911446      NaN  0.0000000      NaN
##  0.11032045  5.835644047  0.0000000      NaN  0.0000000      NaN
##  0.14298028  5.900723701  0.0000000      NaN  0.0000000      NaN
##  0.15333117  0.001237096  0.5832537  0.5727365  0.2724289  0.3689496
##  0.24414545  0.041843375  0.5933157  0.6930946  0.2501094  0.3672583
##  0.24717296  0.005363840  0.5855347  0.5762286  0.2636761  0.3615385
##  0.29186267  0.038180467  0.5940984  0.6801366  0.2507659  0.3661709
##  0.37615343  0.002088687  0.5844254  0.5725664  0.2704595  0.3670794
##  0.57222747  0.041985838  0.5950446  0.7274197  0.2487965  0.3703174
##  0.73301844  2.449870885  0.0000000      NaN  0.0000000      NaN
##  0.77100396  0.191331587  0.5819488      NaN  0.0000000      NaN
##  0.78001638  0.014048926  0.5939276  0.6109304  0.2562363  0.3607058
##  0.96830053  0.005986254  0.5905502  0.5808563  0.2606127  0.3595125
##  0.98979898  1.565802338  0.0000000      NaN  0.0000000      NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.5722275 and lambda
## = 0.04198584.
##
## $'5'$f1_val
## [1] 0.06586826
##
## $'5'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes   22   74
##           no  550 1755
##
##           Accuracy : 0.7401
##           95% CI : (0.7221, 0.7576)
##           No Information Rate : 0.7618
##           P-Value [Acc > NIR] : 0.9937
##
##           Kappa : -0.0028
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.038462
##           Specificity : 0.959541
##           Pos Pred Value : 0.229167
##           Neg Pred Value : 0.761388
##           Prevalence : 0.238234
##           Detection Rate : 0.009163
##           Detection Prevalence : 0.039983
##           Balanced Accuracy : 0.499001
##
##           'Positive' Class : yes
##
##
##

```

```

##
##
##
##
##
## $'1'
## $'1'$model
## Random Forest
##
## 9601 samples
## 6 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7681, 7681, 7681, 7681, 7681, ...
## Resampling results across tuning parameters:
##
## mtry AUC Precision Recall F
## 1 0.8340516 1.0000000 0.01356005 0.06386678
## 3 0.9395373 0.9378831 0.84011543 0.88624213
## 4 0.9462432 0.9343864 0.86701480 0.89937838
## 7 0.7274992 0.9333987 0.91207275 0.92255994
## 8 0.6474083 0.9345026 0.92257365 0.92845568
## 10 0.5594492 0.9337276 0.92607426 0.92983167
## 11 0.5533944 0.9330670 0.92607331 0.92951403
## 12 0.5311951 0.9333159 0.92607379 0.92962965
## 14 0.4890939 0.9302047 0.92585497 0.92796681
## 15 0.4739264 0.9289500 0.92585497 0.92735754
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 10.
##
## $'1'$f1_val
## [1] 0.109834
##
## $'1'$confm
## Confusion Matrix and Statistics
##
## Reference
## Prediction yes no
## yes 43 169
## no 528 1659
##
## Accuracy : 0.7095
## 95% CI : (0.6908, 0.7276)
## No Information Rate : 0.762
## P-Value [Acc > NIR] : 1
##
## Kappa : -0.0219
##
## McNemar's Test P-Value : <2e-16
##

```

```

##          Sensitivity : 0.07531
##          Specificity : 0.90755
##          Pos Pred Value : 0.20283
##          Neg Pred Value : 0.75857
##          Prevalence : 0.23802
##          Detection Rate : 0.01792
##          Detection Prevalence : 0.08837
##          Balanced Accuracy : 0.49143
##
##          'Positive' Class : yes
##
##
##
## $'2'
## $'2'$model
## Random Forest
##
## 9600 samples
##      6 predictor
##      2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7680, 7680, 7680, 7680, 7680, ...
## Resampling results across tuning parameters:
##
##      mtry  AUC          Precision  Recall      F
##      1    0.8302486  1.0000000  0.002625821  0.01731016
##      3    0.9377880  0.9414006  0.834354486  0.88449297
##      4    0.9434182  0.9364035  0.865426696  0.89937172
##      7    0.7928893  0.9369335  0.911597374  0.92398344
##      8    0.7038235  0.9367652  0.921663020  0.92902274
##     10    0.6356841  0.9345435  0.925820569  0.93004998
##     11    0.6125494  0.9333441  0.925164114  0.92909673
##     12    0.6000030  0.9314347  0.924945295  0.92805959
##     14    0.5702973  0.9290084  0.924507659  0.92662005
##     15    0.5584111  0.9283400  0.924070022  0.92607973
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 10.
##
## $'2'$f1_val
## [1] 0.118239
##
## $'2'$confm
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  yes   no
##          yes   47  176
##          no   525 1652
##
##          Accuracy : 0.7079
##          95% CI : (0.6893, 0.7261)

```

```

##      No Information Rate : 0.7617
##      P-Value [Acc > NIR] : 1
##
##      Kappa : -0.0179
##
##      McNemar's Test P-Value : <2e-16
##
##      Sensitivity : 0.08217
##      Specificity : 0.90372
##      Pos Pred Value : 0.21076
##      Neg Pred Value : 0.75884
##      Prevalence : 0.23833
##      Detection Rate : 0.01958
##      Detection Prevalence : 0.09292
##      Balanced Accuracy : 0.49294
##
##      'Positive' Class : yes
##
##
##
## $'3'
## $'3'$model
## Random Forest
##
## 9600 samples
##      6 predictor
##      2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
##      mtry  AUC          Precision  Recall      F
##      1     0.8335961  1.0000000  0.002625821  0.02590872
##      3     0.9355564  0.9414090  0.840330425  0.88786102
##      4     0.9461599  0.9343939  0.864607321  0.89806793
##      7     0.7930380  0.9338374  0.912724910  0.92309770
##      8     0.7239428  0.9340612  0.922347663  0.92809013
##      10    0.6574238  0.9329622  0.927160234  0.92996755
##      11    0.6388138  0.9322938  0.926285439  0.92919648
##      12    0.6149003  0.9303613  0.926723075  0.92844170
##      14    0.5833562  0.9296924  0.926941416  0.92824319
##      15    0.5676825  0.9276230  0.926503301  0.92699225
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 10.
##
## $'3'$f1_val
## [1] 0.1985112
##
## $'3'$confm
## Confusion Matrix and Statistics
##

```

```

##           Reference
## Prediction  yes   no
##           yes    80 154
##           no    491 1675
##
##           Accuracy : 0.7312
##           95% CI : (0.713, 0.7489)
##           No Information Rate : 0.7621
##           P-Value [Acc > NIR] : 0.9998
##
##           Kappa : 0.0701
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.14011
##           Specificity : 0.91580
##           Pos Pred Value : 0.34188
##           Neg Pred Value : 0.77331
##           Prevalence : 0.23792
##           Detection Rate : 0.03333
##           Detection Prevalence : 0.09750
##           Balanced Accuracy : 0.52795
##
##           'Positive' Class : yes
##
##
## $'4'
## $'4'$model
## Random Forest
##
## 9600 samples
##    6 predictor
##    2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
##  mtry  AUC          Precision Recall      F
##    1   0.8298945  1.0000000  0.01159499  0.04495334
##    3   0.9331156  0.9373460  0.83508404  0.88302029
##    4   0.9423735  0.9311381  0.86832771  0.89847954
##    7   0.7982230  0.9277492  0.90747852  0.91739694
##    8   0.7065197  0.9287166  0.91491405  0.92168138
##   10   0.6305422  0.9255611  0.91600910  0.92064919
##   11   0.6072451  0.9264572  0.91469666  0.92042937
##   12   0.5854844  0.9228110  0.91425951  0.91838695
##   14   0.5548390  0.9213188  0.91535312  0.91817190
##   15   0.5464007  0.9202180  0.91425807  0.91708677
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 8.

```

```

##
## $'4'$f1_val
## [1] 0.08994709
##
## $'4'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes   34  152
##           no   537 1677
##
##           Accuracy : 0.7129
##           95% CI : (0.6944, 0.731)
##           No Information Rate : 0.7621
##           P-Value [Acc > NIR] : 1
##
##           Kappa : -0.0307
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.05954
##           Specificity : 0.91689
##           Pos Pred Value : 0.18280
##           Neg Pred Value : 0.75745
##           Prevalence : 0.23792
##           Detection Rate : 0.01417
##           Detection Prevalence : 0.07750
##           Balanced Accuracy : 0.48822
##
##           'Positive' Class : yes
##
##
##
## $'5'
## $'5'$model
## Random Forest
##
## 9599 samples
## 6 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##
## mtry  AUC          Precision Recall      F
## 1      0.8325350      NaN    0.0000000      NaN
## 3      0.9382353  0.9332901  0.8439825  0.8862211
## 4      0.9485859  0.9284373  0.8778993  0.9023586
## 7      0.7914873  0.9334663  0.9164114  0.9248003
## 8      0.7227009  0.9324627  0.9199125  0.9260960
## 10     0.6230053  0.9326191  0.9223195  0.9273996

```



```

## 11 0.6042595 0.9325230 0.9242888 0.9283550
## 12 0.5859090 0.9317753 0.9253829 0.9285325
## 14 0.5573494 0.9302556 0.9238512 0.9270012
## 15 0.5367351 0.9286138 0.9238512 0.9261866
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 12.
##
## $'5'$f1_val
## [1] 0.110971
##
## $'5'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##          yes   44  177
##          no   528 1652
##
##           Accuracy : 0.7064
##           95% CI : (0.6877, 0.7245)
##       No Information Rate : 0.7618
##       P-Value [Acc > NIR] : 1
##
##           Kappa : -0.0252
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.07692
##           Specificity : 0.90323
##       Pos Pred Value : 0.19910
##       Neg Pred Value : 0.75780
##           Prevalence : 0.23823
##       Detection Rate : 0.01833
##       Detection Prevalence : 0.09204
##       Balanced Accuracy : 0.49007
##
##       'Positive' Class : yes
##
##
##
##
##
##
##
##
##
## $'1'
## $'1'$model
## k-Nearest Neighbors
##
## 9601 samples
## 15 predictor
## 2 classes: 'yes', 'no'

```

```

##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7681, 7681, 7681, 7681, 7681, ...
## Resampling results across tuning parameters:
##
##   k   AUC      Precision  Recall    F
##   2  0.1361611  0.8233147  0.8805792  0.8508726
##   3  0.2020149  0.8326205  0.8707285  0.8511398
##   4  0.2417961  0.8305367  0.8543276  0.8421786
##   5  0.2683598  0.8453635  0.8464545  0.8458770
##  10  0.4036427  0.8410745  0.8475471  0.8442668
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
##
## $'1'$f1_val
## [1] 0.3675958
##
## $'1'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes  214  365
##           no   357 1463
##
##           Accuracy : 0.699
##           95% CI : (0.6802, 0.7174)
##           No Information Rate : 0.762
##           P-Value [Acc > NIR] : 1.0000
##
##           Kappa : 0.1743
##
## Mcnemar's Test P-Value : 0.7945
##
##           Sensitivity : 0.3748
##           Specificity : 0.8003
##           Pos Pred Value : 0.3696
##           Neg Pred Value : 0.8038
##           Prevalence : 0.2380
##           Detection Rate : 0.0892
##           Detection Prevalence : 0.2414
##           Balanced Accuracy : 0.5876
##
##           'Positive' Class : yes
##
##
##
## $'2'
## $'2'$model
## k-Nearest Neighbors
##
## 9600 samples

```

```

## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7680, 7680, 7680, 7680, 7680, 7680, ...
## Resampling results across tuning parameters:
##
## k   AUC       Precision Recall    F
## 2   0.1441727 0.8172311 0.8704595 0.8427944
## 3   0.2084413 0.8314812 0.8634573 0.8469245
## 4   0.2498527 0.8352245 0.8560175 0.8453306
## 5   0.2793481 0.8428885 0.8492341 0.8457637
## 10  0.4087602 0.8438838 0.8485777 0.8460734
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
##
## $'2'$f1_val
## [1] 0.4007156
##
## $'2'$confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction yes  no
##          yes 221 330
##          no 351 1498
##
##           Accuracy : 0.7163
##           95% CI : (0.6977, 0.7342)
##       No Information Rate : 0.7617
##       P-Value [Acc > NIR] : 1.0000
##
##           Kappa : 0.2085
##
## Mcnemar's Test P-Value : 0.4434
##
##           Sensitivity : 0.38636
##           Specificity : 0.81947
##       Pos Pred Value : 0.40109
##       Neg Pred Value : 0.81017
##           Prevalence : 0.23833
##       Detection Rate : 0.09208
##       Detection Prevalence : 0.22958
##       Balanced Accuracy : 0.60292
##
##       'Positive' Class : yes
##
##
## $'3'
## $'3'$model
## k-Nearest Neighbors

```

```

##
## 9600 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
## k AUC Precision Recall F
## 2 0.1363338 0.8200137 0.8790426 0.8484538
## 3 0.2078433 0.8324206 0.8722631 0.8518093
## 4 0.2527100 0.8349515 0.8584823 0.8464787
## 5 0.2896776 0.8414663 0.8460144 0.8436588
## 10 0.4180243 0.8501025 0.8348638 0.8423424
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
##
## $'3'$f1_val
## [1] 0.4333895
##
## $'3'$confm
## Confusion Matrix and Statistics
##
##          Reference
## Prediction yes  no
##          yes 261 356
##          no 310 1473
##
##          Accuracy : 0.7225
##          95% CI : (0.7041, 0.7403)
##          No Information Rate : 0.7621
##          P-Value [Acc > NIR] : 1.00000
##
##          Kappa : 0.2554
##
##          McNemar's Test P-Value : 0.08121
##
##          Sensitivity : 0.4571
##          Specificity : 0.8054
##          Pos Pred Value : 0.4230
##          Neg Pred Value : 0.8261
##          Prevalence : 0.2379
##          Detection Rate : 0.1087
##          Detection Prevalence : 0.2571
##          Balanced Accuracy : 0.6312
##
##          'Positive' Class : yes
##
##
## $'4'

```

```

## $'4'$model
## k-Nearest Neighbors
##
## 9600 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7680, 7680, 7680, 7681, 7680, ...
## Resampling results across tuning parameters:
##
## k AUC Precision Recall F
## 2 0.1382653 0.8136408 0.8637297 0.8377682
## 3 0.2066693 0.8325459 0.8661362 0.8488567
## 4 0.2465014 0.8350916 0.8482007 0.8414494
## 5 0.2836692 0.8444190 0.8407614 0.8423856
## 10 0.4175974 0.8415670 0.8492943 0.8451582
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
##
## $'4'$f1_val
## [1] 0.4320675
##
## $'4'$confm
## Confusion Matrix and Statistics
##
## Reference
## Prediction yes no
## yes 261 354
## no 310 1475
##
## Accuracy : 0.7233
## 95% CI : (0.705, 0.7412)
## No Information Rate : 0.7621
## P-Value [Acc > NIR] : 0.99999
##
## Kappa : 0.2567
##
## McNemar's Test P-Value : 0.09517
##
## Sensitivity : 0.4571
## Specificity : 0.8065
## Pos Pred Value : 0.4244
## Neg Pred Value : 0.8263
## Prevalence : 0.2379
## Detection Rate : 0.1087
## Detection Prevalence : 0.2562
## Balanced Accuracy : 0.6318
##
## 'Positive' Class : yes
##
##

```

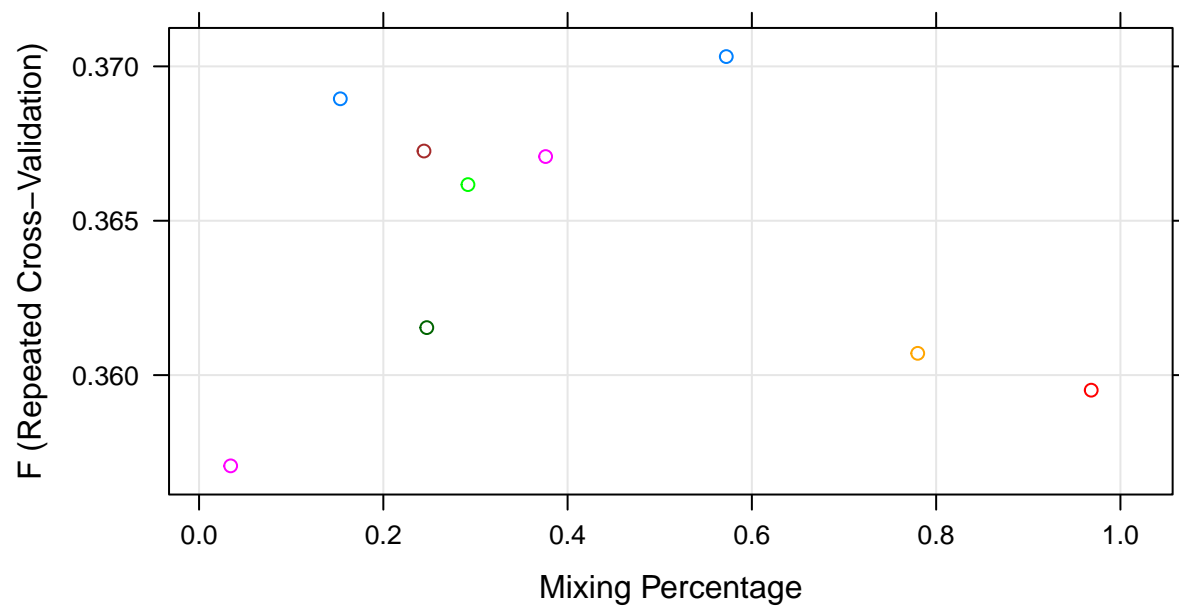

```

##
## $'5'
## $'5'$model
## k-Nearest Neighbors
##
## 9599 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##
## k AUC Precision Recall F
## 2 0.1421366 0.8161043 0.8667396 0.8405315
## 3 0.2004090 0.8382867 0.8669584 0.8522979
## 4 0.2486435 0.8381891 0.8571116 0.8474112
## 5 0.2869835 0.8489578 0.8577681 0.8532106
## 10 0.4402926 0.8390127 0.8557987 0.8471668
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.
##
## $'5'$f1_val
## [1] 0.4092466
##
## $'5'$confm
## Confusion Matrix and Statistics
##
##          Reference
## Prediction yes  no
##          yes 238 364
##          no 334 1465
##
##          Accuracy : 0.7093
##          95% CI : (0.6907, 0.7274)
##          No Information Rate : 0.7618
##          P-Value [Acc > NIR] : 1.0000
##
##          Kappa : 0.2132
##
## Mcnemar's Test P-Value : 0.2724
##
##          Sensitivity : 0.41608
##          Specificity : 0.80098
##          Pos Pred Value : 0.39535
##          Neg Pred Value : 0.81434
##          Prevalence : 0.23823
##          Detection Rate : 0.09913
##          Detection Prevalence : 0.25073
##          Balanced Accuracy : 0.60853
##
##          'Positive' Class : yes

```


```
##  
##  
##  
##  
##  
##
```

```
# Choose best model in terms of F1 score on test dataset  
for (m in models) {  
  f <- get(paste0(m, '_cv'))[[1]]$f1_val  
  n <- 1  
  
  for (i in 2 : n_fold) {  
    f1 <- get(paste0(m, '_cv'))[[i]]$f1_val  
    if (f1 > f) {  
      f <- f1  
      n <- i  
    }  
  }  
  
  assign(paste0(m, '_best'), get(paste0(m, '_cv'))[[i]])  
}  
  
# Record best model for each method  
#results <- as.data.frame(cbind(glmnet_best, svmLinear_best, knn_best, rf_best))  
results <- as.data.frame(cbind(glmnet_best, knn_best, rf_best))  
  
plot(results$glmnet_best$model)
```




```
#plot(results$svmLinear_best$model)
plot(results$knn_best$model)
```





```
plot(results$rf_best$model)
```



```
for (i in 1 : 3) {
  cat(rep('\n', 3))
  print(results[[i]])
  cat(rep('\n', 3))
}
```

```
##
##
##
## $model
## glmnet
##
## 9599 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##
## alpha      lambda      AUC      Precision Recall      F
## 0.03431740 0.089848573 0.5870402 0.8210220 0.2284464 0.3570610
## 0.08249613 0.360573416 0.5911446      NaN 0.0000000      NaN
## 0.11032045 5.835644047 0.0000000      NaN 0.0000000      NaN
## 0.14298028 5.900723701 0.0000000      NaN 0.0000000      NaN
```

```

## 0.15333117 0.001237096 0.5832537 0.5727365 0.2724289 0.3689496
## 0.24414545 0.041843375 0.5933157 0.6930946 0.2501094 0.3672583
## 0.24717296 0.005363840 0.5855347 0.5762286 0.2636761 0.3615385
## 0.29186267 0.038180467 0.5940984 0.6801366 0.2507659 0.3661709
## 0.37615343 0.002088687 0.5844254 0.5725664 0.2704595 0.3670794
## 0.57222747 0.041985838 0.5950446 0.7274197 0.2487965 0.3703174
## 0.73301844 2.449870885 0.0000000 NaN 0.0000000 NaN
## 0.77100396 0.191331587 0.5819488 NaN 0.0000000 NaN
## 0.78001638 0.014048926 0.5939276 0.6109304 0.2562363 0.3607058
## 0.96830053 0.005986254 0.5905502 0.5808563 0.2606127 0.3595125
## 0.98979898 1.565802338 0.0000000 NaN 0.0000000 NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.5722275 and lambda
## = 0.04198584.
##
## $f1_val
## [1] 0.06586826
##
## $confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes   22   74
##           no   550 1755
##
##           Accuracy : 0.7401
##           95% CI : (0.7221, 0.7576)
##           No Information Rate : 0.7618
##           P-Value [Acc > NIR] : 0.9937
##
##           Kappa : -0.0028
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.038462
##           Specificity : 0.959541
##           Pos Pred Value : 0.229167
##           Neg Pred Value : 0.761388
##           Prevalence : 0.238234
##           Detection Rate : 0.009163
##           Detection Prevalence : 0.039983
##           Balanced Accuracy : 0.499001
##
##           'Positive' Class : yes
##
##
##
##
##
##
##
##
##

```

```

## $model
## k-Nearest Neighbors
##
## 9599 samples
## 15 predictor
## 2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##
## k   AUC      Precision Recall      F
## 2   0.1421366 0.8161043 0.8667396 0.8405315
## 3   0.2004090 0.8382867 0.8669584 0.8522979
## 4   0.2486435 0.8381891 0.8571116 0.8474112
## 5   0.2869835 0.8489578 0.8577681 0.8532106
## 10  0.4402926 0.8390127 0.8557987 0.8471668
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.
##
## $f1_val
## [1] 0.4092466
##
## $confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction yes  no
##      yes  238  364
##      no   334 1465
##
##           Accuracy : 0.7093
##           95% CI : (0.6907, 0.7274)
##      No Information Rate : 0.7618
##      P-Value [Acc > NIR] : 1.0000
##
##           Kappa : 0.2132
##
##      McNemar's Test P-Value : 0.2724
##
##           Sensitivity : 0.41608
##           Specificity : 0.80098
##           Pos Pred Value : 0.39535
##           Neg Pred Value : 0.81434
##           Prevalence : 0.23823
##           Detection Rate : 0.09913
##           Detection Prevalence : 0.25073
##           Balanced Accuracy : 0.60853
##
##           'Positive' Class : yes
##
##

```

```

##
##
##
##
##
## $model
## Random Forest
##
## 9599 samples
##    6 predictor
##    2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 7679, 7679, 7679, 7679, 7680, 7679, ...
## Resampling results across tuning parameters:
##
##  mtry  AUC          Precision Recall      F
##    1   0.8325350         NaN 0.0000000      NaN
##    3   0.9382353  0.9332901 0.8439825 0.8862211
##    4   0.9485859  0.9284373 0.8778993 0.9023586
##    7   0.7914873  0.9334663 0.9164114 0.9248003
##    8   0.7227009  0.9324627 0.9199125 0.9260960
##   10   0.6230053  0.9326191 0.9223195 0.9273996
##   11   0.6042595  0.9325230 0.9242888 0.9283550
##   12   0.5859090  0.9317753 0.9253829 0.9285325
##   14   0.5573494  0.9302556 0.9238512 0.9270012
##   15   0.5367351  0.9286138 0.9238512 0.9261866
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 12.
##
## $f1_val
## [1] 0.110971
##
## $confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##          yes   44  177
##          no   528 1652
##
##           Accuracy : 0.7064
##           95% CI : (0.6877, 0.7245)
##        No Information Rate : 0.7618
##        P-Value [Acc > NIR] : 1
##
##           Kappa : -0.0252
##
##  McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.07692

```

```
##          Specificity : 0.90323
##      Pos Pred Value : 0.19910
##      Neg Pred Value : 0.75780
##          Prevalence : 0.23823
##      Detection Rate : 0.01833
## Detection Prevalence : 0.09204
##      Balanced Accuracy : 0.49007
##
##      'Positive' Class : yes
##
##
##
##
```

```
save.image("D:/Yaxin/HKBU BM/Courses/Sem 2/ECON7860 Big Data Analytics for Business (S11)/Group Project,
```