

Modeling_ind.R

yaxin

2021-04-10

```
# Uncomment if packages not installed
## install.packages("psych")
## install.packages("caret")
## install.packages("randomForest")
## install.packages("MLmetrics")
## install.packages("doParallel")
## install.packages("kernlab")
## install.packages("glmnet")

# Load data
setwd('D:\\Yaxin\\HKBU BM\\Courses\\Sem 2\\ECON7860 Big Data Analytics for Business (S11)\\Group Project')
rawData <- read.csv2("HR_comma_sep.csv", sep = ',')

# Transform feature types
transform_feature <- function(X) {
  X$satisfaction_level <- as.numeric(X$satisfaction_level)
  X$last_evaluation <- as.numeric(X$last_evaluation)
  X$Work_accident <- as.factor(X$Work_accident)
  X$promotion_last_5years <- as.factor(X$promotion_last_5years)
  X$sales <- as.factor(X$sales)
  X$salary <- as.factor(X$salary)
  X$left <- factor(ifelse(X$left == 0, 'no', 'yes'), levels = c('yes', 'no'))
  return(X)
}

rawData <- transform_feature(rawData)
summary(rawData)
```

```
## satisfaction_level last_evaluation number_project average_montly_hours
## Min. :0.0900 Min. :0.3600 Min. :2.000 Min. : 96.0
## 1st Qu.:0.4400 1st Qu.:0.5600 1st Qu.:3.000 1st Qu.:156.0
## Median :0.6400 Median :0.7200 Median :4.000 Median :200.0
## Mean :0.6128 Mean :0.7161 Mean :3.803 Mean :201.1
## 3rd Qu.:0.8200 3rd Qu.:0.8700 3rd Qu.:5.000 3rd Qu.:245.0
## Max. :1.0000 Max. :1.0000 Max. :7.000 Max. :310.0
##
## time_spend_company Work_accident left promotion_last_5years
## Min. : 2.000 0:12830 yes: 3571 0:14680
```

```
## 1st Qu.: 3.000      1: 2169      no :11428    1: 319
## Median : 3.000
## Mean   : 3.498
## 3rd Qu.: 4.000
## Max.   :10.000
##
##          sales          salary
## sales      :4140    high :1237
## technical  :2720    low  :7316
## support    :2229    medium:6446
## IT         :1227
## product_mng: 902
## marketing  : 858
## (Other)    :2923
```

```
# Separate target variable
```

```
X <- rawData
y <- X$left
tag <- colnames(X)
tag
```

```
## [1] "satisfaction_level" "last_evaluation" "number_project"
## [4] "average_monthly_hours" "time_spend_company" "Work_accident"
## [7] "left" "promotion_last_5years" "sales"
## [10] "salary"
```

```
# Feature engineering
```

```
## Create dummy variables for "sales" and "salary"
library(psych)
```

```
## Warning: package 'psych' was built under R version 4.0.4
```

```
dummySales <- dummy.code(X$sales)
dummySalary <- dummy.code(X$salary)
colnames(dummySales)
```

```
## [1] "sales" "technical" "support" "IT" "product_mng"
## [6] "marketing" "RandD" "accounting" "hr" "management"
```

```
colnames(dummySalary)
```

```
## [1] "low" "medium" "high"
```

```
### Set "sales" and "low" as the default values respectively
```

```
dummySales <- dummySales[, -c(1)]
dummySalary <- dummySalary[, -c(1)]
```

```
X_dummy <- cbind(X[, -c(9, 10)], dummySales, dummySalary)
tag_dummy <- colnames(X_dummy)
tag_dummy
```

```
## [1] "satisfaction_level" "last_evaluation" "number_project"
## [4] "average_monthly_hours" "time_spend_company" "Work_accident"
## [7] "left" "promotion_last_5years" "technical"
## [10] "support" "IT" "product_mng"
## [13] "marketing" "RandD" "accounting"
## [16] "hr" "management" "medium"
## [19] "high"
```

```
## Create indicator variable for "time_spend_company"
```

```
time_over_5 <- factor(ifelse(X$time_spend_company > 5, 1, 0))
X <- cbind(X[, 1 : 4], time_over_5, X[, 6 : length(X)])
X_dummy <- cbind(X_dummy[, 1 : 4], time_over_5, X_dummy[, 6 : length(X_dummy)])
tag1 <- colnames(X)
tag1
```

```
## [1] "satisfaction_level" "last_evaluation" "number_project"
## [4] "average_monthly_hours" "time_over_5" "Work_accident"
## [7] "left" "promotion_last_5years" "sales"
## [10] "salary"
```

```
# Train(80%)-test(20%)-split (stratified as "left" is unbalanced)
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.0.4
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'ggplot2'
```

```
## The following objects are masked from 'package:psych':
```

```
##
```

```
## %+%, alpha
```

```
## Set seed for replication purpose
```

```
set.seed(7860)
index <- createDataPartition(y, p = 0.8, list = FALSE)
X_train <- X[index, ]
X_test <- cbind(X[-index, 1 : 6], X[-index, 8 : length(X)])
y_test <- X[-index, 'left']
X_dummy_train <- X_dummy[index, ]
X_dummy_test <- cbind(X_dummy[-index, 1 : 6], X_dummy[-index, 8 : length(X_dummy)])
```

```
# Modeling with extracted factors, 5-fold nested CV with random search
```

```
models <- c('svmLinear', 'glmnet', 'rf', 'knn')
```

```
n_cluster <- 10 ## Please set the number of multiprocessing slaves accordingly
```

```

for (m in models) {
  assign(paste0(m, '_best'), list('model' = c(), 'f1_val' = c(),
                                'confm' = c()))

  tune <- 15
  control <- trainControl(method = 'repeatedcv', number = 5, repeats = 2,
                          summaryFunction = prSummary, classProbs = TRUE,
                          search="random", verboseIter = TRUE)

  set.seed(7860)

  require(doParallel)
  cl <- makePSOCKcluster(n_cluster, outfile = '')
  registerDoParallel(cl)

  if (m == 'rf') {
    m1 <- train(left ~ ., data = X_train, method = m,
                metric = 'F', tuneLength = tune, trControl = control)
    rf_best[['model']] <- m1
    rf_best[['f1_val']] <- F_meas(predict(m1, X_test), y_test)
    rf_best[['confm']] <- confusionMatrix(predict(m1, X_test), y_test)
  } else if (m == 'glmnet') {
    m1 <- train(left ~ ., data = cbind(scale(X_dummy_train[ , 1 : 4]), X_dummy_train[ , 5 : length(X_dummy_train)]),
                method = m, family = 'binomial',
                metric = 'F', tuneLength = tune, trControl = control)
    glmnet_best[['model']] <- m1
    glmnet_best[['f1_val']] <- F_meas(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
    glmnet_best[['confm']] <- confusionMatrix(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
  } else if (m == 'knn') {
    m1 <- train(left ~ ., data = cbind(scale(X_dummy_train[ , 1 : 4]), X_dummy_train[ , 5 : length(X_dummy_train)]),
                metric = 'F', tuneLength = tune, trControl = control, tuneGrid = expand.grid(k = c(2, 3, 4, 5, 6, 7, 8, 9, 10, 15, 20, 30, 40, 50)))
    knn_best[['model']] <- m1
    knn_best[['f1_val']] <- F_meas(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
    knn_best[['confm']] <- confusionMatrix(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
  } else {
    m1 <- train(left ~ ., data = cbind(scale(X_dummy_train[ , 1 : 4]), X_dummy_train[ , 5 : length(X_dummy_train)]),
                metric = 'F', tuneLength = tune, trControl = control)
    svmLinear_best[['model']] <- m1
    svmLinear_best[['f1_val']] <- F_meas(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
    svmLinear_best[['confm']] <- confusionMatrix(predict(m1, cbind(scale(X_dummy_test[ , 1 : 4]), X_dummy_test[ , 5 : length(X_dummy_test)])), y_test)
  }

  stopImplicitCluster()
  stopCluster(cl)
}

```

```
## Loading required package: doParallel
```

```
## Warning: package 'doParallel' was built under R version 4.0.4
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 4.0.4
```

```

## Loading required package: iterators

## Warning: package 'iterators' was built under R version 4.0.4

## Loading required package: parallel

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results
## Selecting tuning parameters
## Fitting C = 921 on full training set

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Aggregating results


## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results

## Selecting tuning parameters
## Fitting alpha = 0.153, lambda = 0.00124 on full training set
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 7 on full training set
## Aggregating results
## Selecting tuning parameters
## Fitting k = 3 on full training set


results <- as.data.frame(cbind(glmnet_best, svmLinear_best, knn_best, rf_best))

plot(results$glmnet_best$model)


```




```
plot(results$svmLinear_best$model)
```



```
plot(results$knn_best$model)
```



```
plot(results$rf_best$model)
```

```
for (i in 1 : 4) {
  cat(rep('\n', 3))
  print(results[[i]])
  cat(rep('\n', 3))
}
```

```
##
##
##
## $model
## glmnet
##
## 12000 samples
##    18 predictor
##    2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 9599, 9600, 9600, 9600, 9601, 9600, ...
## Resampling results across tuning parameters:
##
##  alpha      lambda      AUC      Precision  Recall      F
##  0.03431740  0.089848573  0.5978424  0.8278177  0.2070356  0.3309151
##  0.08249613  0.360573416  0.5770630      NaN  0.0000000      NaN
##  0.11032045  5.835644047  0.0000000      NaN  0.0000000      NaN
##  0.14298028  5.900723701  0.0000000      NaN  0.0000000      NaN
```

```

## 0.15333117 0.001237096 0.5764040 0.6574883 0.3579011 0.4631284
## 0.24414545 0.041843375 0.5995441 0.7022834 0.2413432 0.3589979
## 0.24717296 0.005363840 0.5809894 0.6607318 0.3271016 0.4371391
## 0.29186267 0.038180467 0.5992936 0.6919188 0.2430939 0.3595342
## 0.37615343 0.002088687 0.5778859 0.6585455 0.3505511 0.4571495
## 0.57222747 0.041985838 0.5720001 0.7066776 0.2429200 0.3612896
## 0.73301844 2.449870885 0.0000000 NaN 0.0000000 NaN
## 0.77100396 0.191331587 0.5329404 NaN 0.0000000 NaN
## 0.78001638 0.014048926 0.5957931 0.6166058 0.2467702 0.3521881
## 0.96830053 0.005986254 0.5867548 0.6313787 0.2861499 0.3933344
## 0.98979898 1.565802338 0.0000000 NaN 0.0000000 NaN
##
## F was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1533312 and lambda
## = 0.001237096.
##
## $f1_val
## [1] 0.4330275
##
## $confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes  236  140
##           no   478 2145
##
##           Accuracy : 0.7939
##           95% CI : (0.779, 0.8083)
##           No Information Rate : 0.7619
##           P-Value [Acc > NIR] : 1.606e-05
##
##           Kappa : 0.3216
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.33053
##           Specificity : 0.93873
##           Pos Pred Value : 0.62766
##           Neg Pred Value : 0.81777
##           Prevalence : 0.23808
##           Detection Rate : 0.07869
##           Detection Prevalence : 0.12538
##           Balanced Accuracy : 0.63463
##
##           'Positive' Class : yes
##
##
##
##
##
##
##
##
##
##

```

```

## $model
## Support Vector Machines with Linear Kernel
##
## 12000 samples
##    18 predictor
##    2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 9599, 9600, 9600, 9600, 9601, 9600, ...
## Resampling results across tuning parameters:
##
##      C          AUC      Precision  Recall    F
##      0.04464867  0.5978251  0.6159196  0.2726697  0.3771953
##      0.07368128  0.5982951  0.6181018  0.2682988  0.3736204
##      0.09840031  0.5986567  0.6230300  0.2707491  0.3768356
##      0.13818800  0.5987977  0.6198423  0.2703982  0.3759589
##      0.15388959  0.5988409  0.6206067  0.2712751  0.3769227
##      0.39561848  0.5990424  0.6222821  0.2716238  0.3775710
##      0.40826969  0.5990663  0.6219923  0.2717996  0.3776369
##      0.64974256  0.5991235  0.6216020  0.2707497  0.3766250
##      1.56081702  0.5991627  0.6211610  0.2698747  0.3757322
##      11.98711010  0.5992446  0.6218698  0.2691738  0.3751156
##      63.79080326  0.5998981  0.6265558  0.2432994  0.3771364
##      94.68470164  0.5970787  0.6191147  0.2656118  0.3713476
##      103.98601383  0.5992500  0.6204117  0.2698740  0.3755295
##      736.48273800  0.4672218  0.6814965  0.1039190  0.3632910
##      920.95367219  0.5226205  0.6595962  0.1474811  0.4031545
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was C = 920.9537.
##
## $f1_val
## [1] NA
##
## $confm
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes    0    0
##           no   714 2285
##
##           Accuracy : 0.7619
##           95% CI : (0.7463, 0.7771)
##           No Information Rate : 0.7619
##           P-Value [Acc > NIR] : 0.51
##
##           Kappa : 0
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.0000
##           Specificity : 1.0000

```

```

##          Pos Pred Value :    NaN
##          Neg Pred Value : 0.7619
##          Prevalence : 0.2381
##          Detection Rate : 0.0000
##          Detection Prevalence : 0.0000
##          Balanced Accuracy : 0.5000
##
##          'Positive' Class : yes
##
##
##
##
##
##
##
##
##
##
## $model
## k-Nearest Neighbors
##
## 12000 samples
##    18 predictor
##    2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 9599, 9600, 9600, 9600, 9601, 9600, ...
## Resampling results across tuning parameters:
##
##  k  AUC          Precision Recall      F
##  2  0.1149322  0.8069732  0.8909706  0.8467571
##  3  0.1727101  0.8086860  0.8923674  0.8483781
##  4  0.2039229  0.8087259  0.8811679  0.8433248
##  5  0.2307771  0.8194162  0.8783701  0.8477608
## 10  0.3251063  0.8170888  0.8689163  0.8420971
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
##
## $f1_val
## [1] 0.8586456
##
## $confm
## Confusion Matrix and Statistics
##
##          Reference
## Prediction yes  no
##          yes  654 154
##          no   60 2131
##
##          Accuracy : 0.9286
##          95% CI : (0.9188, 0.9376)
##          No Information Rate : 0.7619
##          P-Value [Acc > NIR] : < 2.2e-16
##

```

```

##                Kappa : 0.8118
##
## Mcnemar's Test P-Value : 2.053e-10
##
##                Sensitivity : 0.9160
##                Specificity : 0.9326
##                Pos Pred Value : 0.8094
##                Neg Pred Value : 0.9726
##                Prevalence : 0.2381
##                Detection Rate : 0.2181
##                Detection Prevalence : 0.2694
##                Balanced Accuracy : 0.9243
##
##                'Positive' Class : yes
##
##
##
##
##
##
##
##
##
##
## $model
## Random Forest
##
## 12000 samples
##      9 predictor
##      2 classes: 'yes', 'no'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 9599, 9600, 9600, 9600, 9601, 9600, ...
## Resampling results across tuning parameters:
##
##  mtry  AUC          Precision Recall      F
##    1   0.8450443  0.9977625  0.1322784  0.2209106
##    2   0.9509879  0.9855131  0.7045902  0.8214880
##    3   0.9660519  0.9647894  0.8781891  0.9193528
##    4   0.8676781  0.9560223  0.9163420  0.9357037
##    7   0.4985774  0.9569444  0.9432960  0.9500074
##    8   0.4618902  0.9554648  0.9436457  0.9494454
##   10   0.4150500  0.9531331  0.9448731  0.9489268
##   12   0.3824791  0.9517627  0.9441723  0.9478962
##   16   0.3468032  0.9479437  0.9443465  0.9460821
##   17   0.3407206  0.9459145  0.9441714  0.9449895
##   18   0.3365850  0.9464936  0.9448716  0.9456151
##
## F was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 7.
##
## $f1_val
## [1] 0.9579832
##
## $confm

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  yes   no
##           yes 684   30
##           no  30 2255
##
##           Accuracy : 0.98
##           95% CI : (0.9743, 0.9847)
##           No Information Rate : 0.7619
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.9449
##
## Mcnemar's Test P-Value : 1
##
##           Sensitivity : 0.9580
##           Specificity : 0.9869
##           Pos Pred Value : 0.9580
##           Neg Pred Value : 0.9869
##           Prevalence : 0.2381
##           Detection Rate : 0.2281
##           Detection Prevalence : 0.2381
##           Balanced Accuracy : 0.9724
##
##           'Positive' Class : yes
##
##
##
##
##

```

```

#save.image("D:/Yaxin/HKBU BM/Courses/Sem 2/ECON7860 Big Data Analytics for Business (S11)/Group Project")

```